

Realizing executable programs from an abstraction of the policy

Senne Deproost
senne.deproost@vub.be

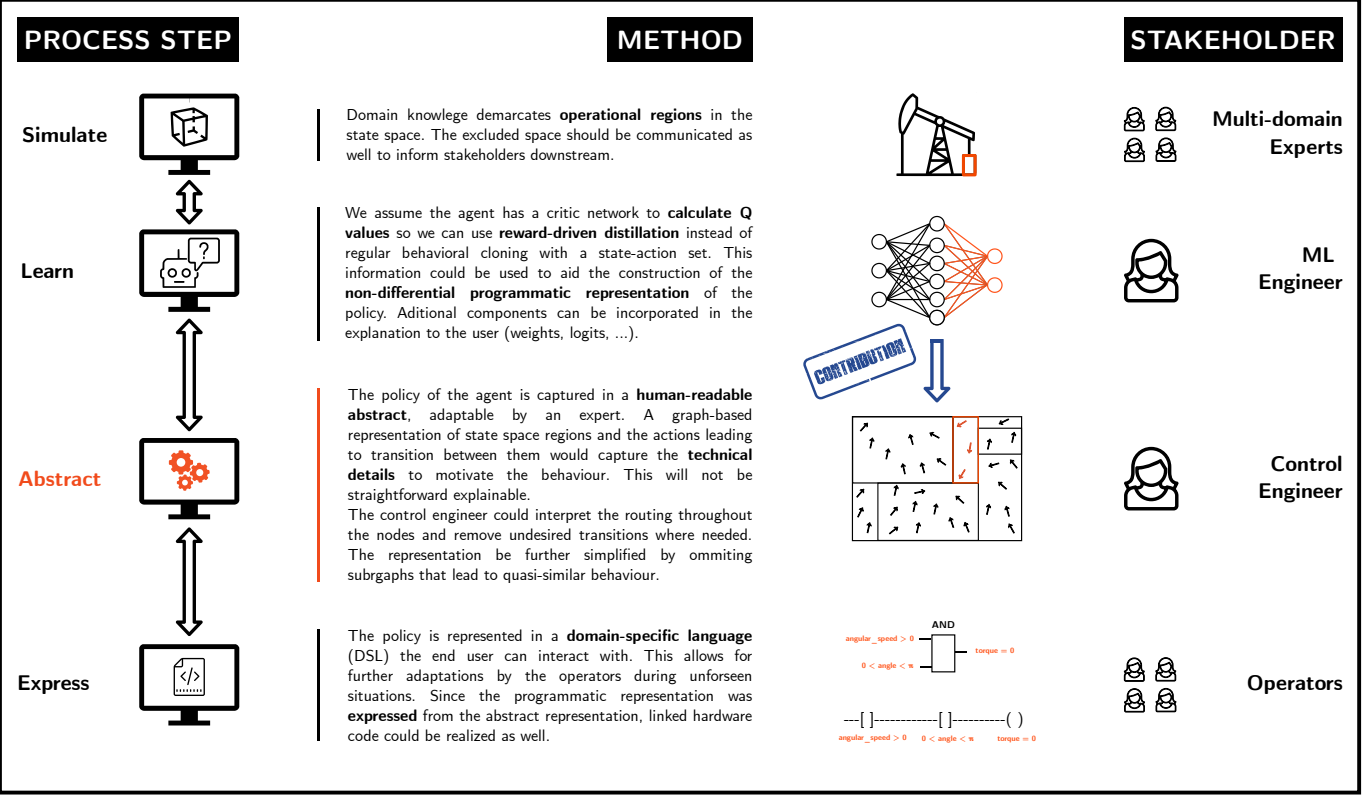
Ann Nowé
ann.nowe@vub.be

Approaches in literature



- **Template-based parametric search [1]**
 - RL policy guides search towards good parameters within the program
 - Structure of the program is fixed
- **Program latent space [2]**
 - Encode embeddings with variational autoencoder so new programs can be generated
 - Requires set of program execution traces
- **Syntax tree via genetic algorithm [3]**
 - Using evolutionary algorithm to generate the abstract syntax tree of the program
 - Slow convergence

Open gaps

- Automatic state space demarcation
- Deepening statements based on expression-simplicity tradeoff
- Guarantees of safe operation
 - Executability guarantee
- Different target languages



Slider crank validation



The slider crank setup is a mechanical system where we translate a linear motion into a rotational one. The difficulty comes from the changing goal angular speed as well as the introduction of other forces (e.g. a spring).

To find a good policy, the **demarcation** of the state space should be more detailed in regions where change in behaviour is more crucial. Within, a programmatic policy would need to be more expressive to better mimic a trained RL agent.

Sources

[1] A. Venna, V. Marai, R. Singh, P. Kóhli, and S. Chaudhuri, 'Programmatically Interpretable Reinforcement Learning', in Proceedings of the 35th International Conference on Machine Learning, PMLR, Jul. 2018, pp. 5045-5054.
[2] D. Hein, S. Udluft, and T. A. Runkler, 'Interpretable policies for reinforcement learning by genetic programming', Engineering Applications of Artificial Intelligence, vol. 76, pp. 158-169, Nov. 2018.
[3] D. Trivedi, J. Zhang, S.-H. Sun, and J. J. Lim, 'Learning to Synthesize Programs as Interpretable and Generalizable Policies', arXiv, Jan. 31, 2022.